

Introduction

...This capability allows readers to quickly establish an initial understanding without reading every passage in detail. However, fluent summaries can also introduce a less visible problem: different AI-generated abstractions of the same document may foreground different arguments, contributions, or implications.

For example, when several LLMs summarize the same research paper, one may emphasize the evaluated system, another the observed user behavior, and another the broader design implications. None of these interpretations is necessarily factually incorrect, yet each directs the reader along a different interpretive path [1, 25, 36]. The challenge is therefore not only whether an AI-generated statement is correct, but whether readers can recover how a particular abstraction relates to the source material from which it was constructed [16, 22].

Existing AI-assisted reading systems have substantially improved how readers access and inspect long documents [2, 12, 20]...

Extract source text



Structure content



Set conversation context



A



Check reading progress

While prompting LLMs can appear effortless, designing effective prompting strategies requires identifying the contexts in which these LLMs' errors arise, devising prompting strategies to overcome them, and systematically assessing those strategies' effectiveness.

Tasks typically handled by 'prompt engineers':

These tasks fall on so-called "prompt engineers"—the designers, domain experts, and any other end-user or professional attempting to improve an LLM's output—and are challenging tasks even for LLM experts.

While prompting LLMs can appear effortless, designing effective prompting strategies requires identifying the contexts in which these LLMs' errors arise, devising prompting strategies to overcome them, and systematically assessing those strategies' effectiveness.



Know what to ask

Same paper, different summaries



B



Get the gist

Same paper, different emphases

Different emphases shape different interpretive paths

Different does not necessarily mean incorrect

The key issue is recoverability

C



See it in context

For example, when several LLMs summarize the same research paper, one may emphasize the evaluated system

another the observed user behavior, and another the broader design implications.

The challenge is therefore not only whether an AI-generated statement is correct, but whether readers can recover how a particular abstraction relates to the source material from which it was constructed [16, 22].

None of these interpretations is necessarily factually incorrect, yet each directs the reader along a different interpretive path [1, 25, 36].

Edit

Question

Add Link

Can you give me an example?

D

Example: Why Johnny Can't Prompt (Zamfirescu-Pereira et al.)
System: BotDesigner as a no-code chatbot design tool
Behavior: Non-experts struggle to engineer prompts effectively
Implication: Conversational habits shape prompt design